



powered by Gemini Deep Research, organized with Notebook LM, developed with chatGPT

produced by [Nicole Dickens, Fractional AI Consultant](#)

October 04 - October 10, 2025

Executive Summary: The Industrialization of AI Compute and Intelligence

The week of October 4–10, 2025, marked a critical pivot point for the Artificial Intelligence sector, confirming its full transition from a speculative research field into an industrial-scale manufacturing and deployment domain. The industry's defining conflict has shifted from model capability breakthroughs to securing compute capacity and demonstrating immediate, quantifiable Return on Investment (ROI) at the enterprise level.

A landmark development in the compute wars was the announcement of the strategic partnership between AMD and OpenAI. This agreement secures 6 gigawatts of AMD Instinct GPUs for OpenAI's next-generation infrastructure, institutionalizing the shift toward a dual-source compute backbone and directly challenging NVIDIA's supply chain dominance. This strategic maneuvering reflects the existential need for guaranteed supply to support infrastructure plans that may cost trillions of dollars. By tying equity vesting (up to 160 million shares of AMD common stock) to deployment volumes, the deal serves as a long-term insurance policy for OpenAI, mitigating the systemic vendor lock-in risk associated with relying on a single major hardware supplier for future scaling.

Complementing the infrastructure push, the enterprise sector received a crucial validation metric this week. The IBM-Anthropic collaboration, integrating the Claude large language model into IBM's software portfolio, yielded a verifiable, high-impact benchmark: early adopters reported an average 45% productivity gain for software developers using the new AI-first integrated development environment (IDE). This result transforms Generative AI integration from an experimental cost center into a strategic capital expenditure justified by measurable efficiency and cost savings.

Geopolitically, the world order continues to bifurcate. The European Union (EU) established its competitive position with the announcement of the "Apply AI Strategy," prioritizing strategic

sovereignty and industrial application in key sectors such as healthcare and defense. This industrial policy, backed by the expansion of the "AI Factories" network, stands in contrast to the United States' continued focus on deregulation and rapid infrastructure buildout through accelerated federal permitting.

Key Takeaways for Small and Medium-Sized Businesses (SMBs)

The current environment necessitates rapid AI adoption for Small and Medium-Sized Businesses (SMBs), driven by clear competitive necessity and proven ROI. The SME segment is projected to register the highest Compound Annual Growth Rate (CAGR) of 32.10% in the forecast period, reflecting high demand for accessible AI tools.

Quantifiable Value and Accessible Tools

Generative AI is moving beyond simple text generation to deliver immediate financial results. Data indicates that 51% of SMBs that have adopted Generative AI reported a revenue increase of 10% or more from their efforts, confirming that AI is a value-generating tool rather than just an efficiency lever.

This value is accessible through the democratization of AI via specialized application-layer tools that abstract away the complexity of foundational models. For instance, workflow automation platforms like Zapier AI connect over 7,000 apps to streamline processes such as lead management and email marketing, providing enterprise-level workflow management without requiring specialized coding expertise. In the e-commerce and marketing sphere, tools like GoDaddy's AI Website Builder and Omnisend's AI segment builder offer 24/7 customer support, optimize email send times, and identify customer segments, reducing time-to-market and enhancing engagement. Furthermore, models like ChatGPT Plus enable small teams to upload and analyze complex documents, spreadsheets, and raw data, delivering key insights that save hours of manual work.

The investment surge into AI infrastructure and optimization stacks, such as Modular's \$250 million Series C, is rapidly driving down the deployment cost of advanced AI. Although frontier models are computationally intense, the capability-to-price ratio continues to double reliably every six to eight months. This cost deflation is crucial, as it ensures that third-party vendors can offer high-performance AI services to SMBs at affordable subscription prices, such as \$20 per month for advanced tiers, effectively enabling these small businesses to leverage advanced reasoning power previously reserved for large enterprises.

The Agentic Commerce Pathway

A notable trend impacting SMBs is the strategic vertical integration seen in merger and acquisition (M&A) activities, signaling the foundation for sophisticated Agentic Commerce platforms. Rezolve AI's recently completed foundational acquisitions of Subsquid (decentralized data layer) and Smartpay (digital payment rail) demonstrate the imperative to build end-to-end autonomous commerce architecture. This vertical integration, combining data, intelligence (via its proprietary LLM, brainpowa), and transaction settlement, is designed to enable autonomous commerce. Soon, these integrated platforms will be packaged and delivered to SMBs, allowing

them to automate entire sales, marketing, and customer journeys within a single, unified ecosystem, effectively bridging the human resource gap.

Despite the accessibility of tools, the talent gap remains the most significant long-term structural risk for small firms. Specialized AI skills command a substantial 56% wage premium. Since SMBs cannot typically compete with large enterprises for this prohibitively expensive talent, their long-term strategy must prioritize vendor-managed automation and the internal upskilling of existing staff, rather than relying on external hiring, to navigate the adoption obstacle cited by industry leaders.

Global AI Policy and Governance: Competition and Control

Geopolitical competition around AI capability and governance intensified this week, characterized by divergent strategic pathways in major global blocs: industrial control in Europe versus market acceleration in the United States.

The European Union: Sovereignty and Industrial Application

The European Union moved decisively toward realizing its ambitions of strategic autonomy. The announcement of the "Apply AI Strategy" (October 5) and related Commission Communication (October 8) outlines a definitive shift toward promoting European-made AI tools across key national strategic areas, specifically citing healthcare, defense, and manufacturing. This initiative is explicitly designed to reduce dependency on powerful US and Chinese AI platforms.

This ambition is being physically realized through the expansion of the "AI Factories" network, bolstering the EU's "AI Continent" initiative, which was accompanied by the announcement of €204 million in funding aimed at boosting innovation in digitalization, digital skills, health, and public and automotive sectors. This approach represents a mature realization that regulatory adherence alone does not guarantee economic success; rather, the EU is pairing its stringent regulatory environment (the AI Act) with state-sponsored industrial policy. By simultaneously subsidizing domestic alternatives and enforcing high compliance standards for foreign firms, the EU seeks to use its regulatory high-bar as a competitive advantage to protect market share in high-value sectors.

United States: Deregulation and Infrastructure Acceleration

The US federal government, guided by the "America's AI Action Plan," continues to prioritize rapid technological leadership and domestic job creation. Policy moves focus heavily on streamlining deployment and procurement, including accelerating federal permitting for data center infrastructure and eliminating barriers for federal AI use. This pro-growth, deregulatory stance ensures the fastest possible compute buildout. Alongside infrastructure, the Administration maintains a critical focus on investment in Biosecurity, recognizing the potential for AI to unlock new cures while simultaneously mitigating the dual-use catastrophic risks inherent in rapid technological advancement.

State-Level Fragmentation and Regulatory Depth

The US federal government's focus on speed and deregulation creates inevitable regulatory gaps. In response, all 50 US states, plus territories, have introduced AI-related legislation in the 2025 legislative session, with 38 states adopting or enacting approximately 100 measures this year. This widespread velocity of legislation demonstrates bottom-up regulatory pressure, forcing companies to address specific, concrete risks that the federal government has chosen to bypass.

For example, Arkansas enacted legislation clarifying Intellectual Property (IP) ownership of AI-generated content, assigning ownership to either the person who provides the input data to train the model or the employer, if the content is generated as part of employment duties. Separately, Montana's new "Right to Compute" law mandates that AI systems controlling critical infrastructure must develop risk management policies aligned with established standards, such as the NIST AI Risk Management Framework. For companies operating nationally, compliance now involves a complex 50-state calculus, making nationwide deployment riskier than operating solely within the EU's single, codified AI Act framework.

The following table summarizes the divergent strategic focus areas adopted by key global jurisdictions this week:

Table 1: Key Global AI Policy Moves: Comparison of Jurisdictional Focus (Week of 10/4/2025)

Jurisdiction	Primary Policy Action (Week of Oct 4-10)	Strategic Goal	Mechanism/Instrument
European Union (EU)	"Apply AI Strategy" & AI Factories Expansion	Technological sovereignty, boosting European industrial adoption (defense, healthcare)	New Commission Communication, €204M funding for "AI Continent" ambition
United States (Federal)	Infrastructure Acceleration & Biosecurity Investment	Enhancing US domestic leadership, rapid compute buildout	Accelerating Federal Permitting for Data Centers, Executive Order Biosecurity Focus
US States (E.g., Arkansas/Montana)	IP Rights & Critical Infrastructure Safety	Clarifying legal risks and ownership rights locally	Enacted laws on AI content ownership and mandated NIST-aligned risk management for critical systems

AI Industry Investment: Mega-Deals and Market Friction

The investment landscape is characterized by extreme polarization, massive speculative infrastructure commitments, and a sharp financial focus on the application layer capable of demonstrating immediate monetization. More than 50% of global Venture Capital (VC) funding in 2025 has been directed toward AI.

The Compute Infrastructure Arms Race

The announcement of the 6-gigawatt AMD and OpenAI partnership is the quarter's most consequential hardware deal, solidifying the trend of massive, long-term capital commitments to compute infrastructure. This agreement, which begins with a 1-gigawatt deployment of the AMD Instinct MI450 series in the second half of 2026, aims to power OpenAI's next-generation infrastructure. The unique deal structure involves AMD issuing OpenAI a warrant for up to 160 million shares of AMD common stock. This vesting is tied not only to the scale of deployment (up to 6 GW) but also to AMD achieving certain share-price targets and OpenAI achieving the technical and commercial milestones required for deployment at scale. This mechanism effectively guarantees volume purchases and aligns the long-term financial success of both companies.

The scale of this infrastructure buildout, driven by projects like OpenAI's "Stargate" plan and Meta's commitment to hundreds of billions in data centers, has generated financial friction. Analysts have been warning that the current investment craze, fueled by debt and a fear of missing out (FOMO), mirrors a speculative bubble. This high-stakes environment is further highlighted by the immense burn rates of leading labs; OpenAI, for instance, is projected to burn \$115 billion by 2029.

Venture Capital Focus on Applied AI and Enterprise Infrastructure

Investment capital continues to show clear signs of maturity by bypassing generalized Large Language Models (LLMs) in favor of solutions that demonstrate immediate traction in enterprise adoption and workflow integration. The market has largely priced in the existence of powerful foundational models, and new investment scarcity is focused squarely on monetization pathways, demanding clear, defensible moats based on proprietary data and enterprise workflow.

This focus on monetization pathways is evident in the week's largest announced funding rounds:

- **Vertical Software Dominance:** Judi Health (\$400 million) in health benefits software and Filevine (\$400 million) in legal practice management secured massive rounds. These investments confirm investor appetite for companies that leverage foundational models to solve high-value, domain-specific problems within highly regulated verticals.
- **Inference Optimization:** Modular, the developer of an enterprise AI inference stack, raised \$250 million in Series C financing at a \$1.6 billion valuation. This investment recognizes the crucial value of the optimization layer required to deploy massive models efficiently and affordably in enterprise environments, ensuring deployment efficiency is prioritized above raw hardware power.
- **Agentic Architecture M&A:** Rezolve AI's acquisitions of Subsquid and Smartpay for a low cash consideration plus stock confirmed the M&A trend toward building highly integrated "Agentic Commerce" architecture. This demonstrates the imperative to control the entire workflow—from the data layer and the intelligence layer to transaction settlement—for

high-margin autonomous applications.

Financial System Vulnerability

Institutional monitoring of the AI supply chain intensified this week. On October 10, 2025, the Financial Stability Board (FSB) issued a report that directly addressed vulnerabilities in the AI supply chain, urging financial authorities to monitor the critical reliance on a few third-party service providers. This institutional warning confirms that the technological supply chain risk—hyperscaler and foundational model dependency—has escalated into a recognized systemic financial risk. Regulators fear that a major disruption or failure of one large, concentrated AI infrastructure provider could instantaneously cascade across the regulated banking and financial system.

Table 2: AI Industry Investment Snapshot (10/4/2025 – 10/10/2025)

Company	Focus Sector	Deal Value (Approx.)	Key Strategic Significance
Judi Health	Health Benefits Software	\$400M (Series F + Add-on)	Scaling enterprise AI in highly regulated vertical
Filevine	Legal Practice Management	\$400M (Undisclosed Rounds)	Validation of vertical agentic workflows for complex legal tasks
Modular	AI Inference Stack / Infrastructure	\$250M (Series C, \$1.6B Valuation)	Investment in deployment efficiency (optimization layer) above hardware
AppZen	FinTech / Agentic AI	\$180M (Series D)	Autonomous agents targeting 50%+ replacement of manual finance work
Rezolve Ai (Acquisitions)	Agentic Commerce Infrastructure	~\$10M Cash + 1M Shares	Strategic vertical integration (data/payments) for Agentic Commerce foundation

Breakthroughs in AI Technology: The Rise of Reasoning and Physical AI

The Shift to Real Reasoning

The technological frontier is currently defined by sophistication in planning and verification, shifting the industry focus from achieving raw statistical performance to enabling true functional intelligence. The "State of AI Report 2025," published October 9, confirms that 2025 was the year "reasoning got real." Flagship models from major labs—including OpenAI, Google, Anthropic, and DeepSeek—are now reliably capable of planning, verifying results, self-correction, and operating over extended time horizons to accomplish complex, multi-step tasks.

While OpenAI retains a narrow lead at the frontier, global competition has intensified dramatically. China’s models, specifically DeepSeek, Qwen, and Kimi, have closed the capability gap with GPT-5 across critical benchmarks like reasoning and coding. These models have secured China a credible claim to second place in global capability, validating the effectiveness of sustained, nationally directed investment in compute and model development despite US

export controls, and providing global entities with powerful non-US technology alternatives.

Verified Enterprise Productivity

The IBM and Anthropic partnership yielded the week's most crucial, quantifiable ROI metric. By integrating Anthropic's Claude into IBM's software portfolio, beginning with an AI-first integrated development environment (IDE), the collaboration is specifically designed for enterprise-grade security and governance. Early internal testers at IBM reported productivity gains averaging 45% in tasks such as application modernization, code refactoring, and general software development lifecycle execution. This publicly validated 45% lift in software development productivity sets a new minimum expectation for enterprise AI ROI in 2026, forcing rival LLM providers and systems integrators to structure their offerings around guaranteed, audited outcomes rather than simple API access, a necessity for justifying multi-million dollar software investments.

The Emergence of World Models

The industry is seeing an infrastructural shift toward Embodied AI and Physical AI. This requires the development of "World Models," which are learned simulations of an agent's environment that allow the agent to "imagine" and explore possible sequences of actions. NVIDIA's Cosmos World Foundation Model (WFM) platform continues to evolve, providing the necessary tools for creating, training, and deploying WFMs, positioning the company not only as a hardware supplier but as a foundational ecosystem provider for advanced robotics and autonomous systems. Key competitive developments include Google DeepMind's DreamerV3 and Genie 2, and Meta's Navigation World Model (NWM).

Constraints on Scientific AI Adoption

Despite technical triumphs like AlphaFold, the impact of AI on scientific discovery remains concentrated. A recent report posted to the arXiv repository highlights that the primary barriers to adoption are now social and institutional, rather than technical. These constraints include widespread data fragmentation, significant infrastructure inequities, and cultural gaps between domain experts and machine learning researchers. This analysis calls for reframing AI for science as a collective social project, where cross-disciplinary education and accessible infrastructure are prerequisites for technical progress.

Societal and Economic Implications

The ongoing maturation of AI is reinforcing a deeply polarized economic landscape, creating massive value for the skilled few while accelerating job displacement for routine tasks, alongside mounting regulatory scrutiny on systemic stability.

Labor Market Polarization: The Skill Acceleration Cliff

The labor market is experiencing intense distortion. Data confirms that the rising adoption of generative AI technology accounted for over 10,000 job cuts in the US in the first seven months of 2025, solidifying AI as a primary factor contributing to displacement. In sharp contrast, workers possessing specific AI skills, such as prompt engineering and AI

auditing, now command a monumental 56% wage premium over peers in the same occupation—a rate that more than doubled over the previous year. This extreme wage premium is not merely an indicator of healthy market demand, but rather a reflection of a severe structural bottleneck: the inability of current educational and corporate upskilling systems to keep pace with the accelerating demand for expertise. The pace of necessary skill change in AI-exposed jobs is accelerating 66% faster than last year , prompting business leaders to acknowledge they are entering "one of the largest change management exercises in history". This distortion ensures that the economic value created by AI (the revenue boost seen by SMBs) is captured by a very small, expensive cohort, increasing social and economic pressure across the broader workforce.

Table 3: Benchmarking Enterprise Productivity Gains and Skills Premium

Metric	Figure (2025)	Context/Implication
Enterprise Development Productivity Gain	Average 45% increase	Verified ROI metric from IBM/Anthropic AI-first IDE integration, validating enterprise deployment in SDLC
Wage Premium for AI Skills	56% Higher	Economic indicator of severe talent gap and high value placed on prompt engineering/AI auditing skills
Speed of Skill Change in AI-Exposed Jobs	66% Faster	Accelerating obsolescence of existing skill sets, demanding rapid corporate reskilling programs
AI-Attributable US Job Cuts (YTD 2025)	Over 10,000	Confirms AI's polarizing impact on the labor market (displacement vs. massive value creation)

Systemic Risk and Malicious Use Mitigation

The focus of AI safety has shifted firmly from theoretical long-term existential risk toward concrete, measurable, and immediate harm reduction. This alignment is evident in both regulatory warnings and corporate action.

The Financial Stability Board’s report on concentration risk highlights that the compute supply chain poses an immediate, systemic vulnerability to global financial stability, forcing financial institutions to enhance monitoring and risk mitigation. On the corporate safety front, OpenAI’s October 7, 2025, report detailed the disruption of over 40 malicious networks. These disrupted activities included uses by authoritarian regimes (control and coercion), scams, malicious cyber activity, and covert influence operations. Crucially, the analysis suggests that threat actors are generally "bolting AI onto old playbooks to move faster," rather than using current models to enable fundamentally novel offensive capabilities. This operational finding reinforces the current policy focus on Biosecurity investment and state-level critical infrastructure safety , indicating that resources are being diverted toward mitigating the amplification of known, tangible harms. Furthermore, legal and compliance experts are prioritizing the establishment of

internal governance to eliminate bias and ensure the fair and equitable use of AI systems, reflecting heightened public and regulatory expectations for ethical deployment.

Conclusions and Outlook

The week of October 4–10, 2025, confirms that the AI industry is operating at a new level of industrial maturity, defined by strategic supply chain control and validated monetization metrics. The agreement between AMD and OpenAI demonstrates the high-stakes financial and logistical engineering required to sustain trillion-dollar infrastructure projections, ensuring competition and supply resilience in the face of escalating bubble warnings.

For the enterprise, the IBM/Anthropic 45% productivity benchmark sets a high, measurable standard for ROI, signaling that generalized models are insufficient; the market now requires integrated, verifiable workflow solutions. This drive toward vertical integration is also evident in M&A, where control over the full agentic commerce stack (data, payments, intelligence) is the new competitive imperative.

The most critical long-term macroeconomic challenge remains the highly polarized labor market. While AI adoption is creating massive, immediate value (51% of SMBs seeing 10%+ revenue growth), the 56% wage premium for specialized skills represents a severe structural bottleneck. Sustained competitive advantage in the coming year will depend less on technical innovation and more on the ability of organizations—both large enterprises and SMBs—to rapidly bridge the talent gap through vendor automation and comprehensive internal upskilling programs. Failure to address this widening skill chasm risks concentrating AI's economic benefits among a few while accelerating social displacement across the many.

Sources

Works cited

1. Welcome to State of AI Report 2025, <https://www.stateof.ai/2025-report-launch>
2. AI Deals in 2025: Key Trends in M&A, Private Equity, and Venture Capital - Morgan Lewis, <https://www.morganlewis.com/pubs/2025/09/ai-deals-in-2025-key-trends-in-ma-private-equity-and-venture-capital>
3. AMD and OpenAI announce strategic partnership to deploy 6 ..., <https://openai.com/index/openai-amd-strategic-partnership/>
4. AI gold rush: Why experts fear a massive trillion-dollar crash could be coming, <https://m.economictimes.com/news/international/us/ai-gold-rush-why-experts-fear-a-massive-trillion-dollar-crash-could-be-coming/articleshow/124341802.cms>
5. IBM and Anthropic Partner to Advance Enterprise Software Development with Proven Security and Governance, <https://newsroom.ibm.com/2025-10-07-2025-ibm-and-anthropic-partner-to-advance-enterprise-software-development-with-proven-security-and-governance>
6. IBM stock surges following Anthropic partnership - The American Bazaar, <https://americanbazaaronline.com/2025/10/08/ibm-stock-surges-following-anthropic-partnership-468525/>
7. IBM and Anthropic team up to boost enterprise software development - eeNews Europe, <https://www.eenewseurope.com/en/ibm-and-anthropic-team-up-to-boost-enterprise-software-d>

evelopment/

8. EU to unveil new AI strategy to reduce dependence on US and China: FT, <https://m.economictimes.com/tech/artificial-intelligence/eu-to-unveil-new-ai-strategy-to-reduce-dependence-on-us-and-china-ft/articleshow/124316211.cms>
9. AI Act | Shaping Europe's digital future, <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>
10. AI.Gov | President Trump's AI Strategy and Action Plan, <https://www.ai.gov/>
11. Artificial Intelligence [AI] Market Size, Growth & Trends by 2032 - Fortune Business Insights, <https://www.fortunebusinessinsights.com/industry-reports/artificial-intelligence-market-100114>
12. 2025: the year companies prepare to disrupt how work gets done | World Economic Forum, <https://www.weforum.org/stories/2025/01/ai-2025-workplace/>
13. 12 Best AI Tools for Small Businesses in 2025 - Nextiva, <https://www.nextiva.com/blog/best-ai-tools-for-small-business.html>
14. Top 10 AI Tools for Small Businesses in 2025: Boosting Efficiency Without Replacing Humans - Twilead, <https://twilead.com/top-10-ai-tools-for-small-businesses/>
15. I tried 70+ best AI tools in 2025 - TechRadar, <https://www.techradar.com/best/best-ai-tools>
16. Best AI Tools for Small Business in 2025 (Free & Paid) - Omnisend, <https://www.omnisend.com/blog/ai-tools-for-small-business/>
17. The 45 Best AI Tools in 2025 (Tried & Tested) - Synthesia, <https://www.synthesia.io/post/ai-tools>
18. The Week's 10 Biggest Funding Rounds: Health And AI Lead For Large Financings, <https://news.crunchbase.com/venture/biggest-funding-rounds-health-ai-judi-filevine/>
19. Rezolve Ai Confirms Minimal Dilution from Foundational Acquisitions, <https://natlawreview.com/press-releases/rezolve-ai-confirms-minimal-dilution-foundational-acquisitions>
20. Rezolve Ai Confirms Minimal Dilution from Foundational Acquisitions - GlobeNewswire, <https://www.globenewswire.com/news-release/2025/10/10/3164822/0/en/Rezolve-Ai-Confirms-Minimal-Dilution-from-Foundational-Acquisitions.html>
21. The Fearless Future: 2025 Global AI Jobs Barometer - PwC, <https://www.pwc.com/gx/en/issues/artificial-intelligence/ai-jobs-barometer.html>
22. National Grid Partners unveils results – and investment in Aina Climate AI Ventures – at 2025 NextGrid Alliance Summit, <https://www.nationalgrid.com/media-centre/press-releases/national-grid-announces-investment-in-Aina-climate-AI-ventures>
23. America's AI Action Plan - The White House, <https://www.whitehouse.gov/wp-content/uploads/2025/07/Americas-AI-Action-Plan.pdf>
24. Artificial Intelligence 2025 Legislation - National Conference of State Legislatures, <https://www.ncsl.org/technology-and-communication/artificial-intelligence-2025-legislation>
25. FSB issues fourth AI report | Global Regulation Tomorrow, <https://www.regulationtomorrow.com/africa/fintech-africa/fsb-issues-fourth-ai-report/>
26. 5+ Notable Models of 2025 that You Should Know (AI 101 Guide Recap) - Turing Post, <https://www.turingpost.com/p/models-recap-2025>
27. AI for Scientific Discovery is a Social Problem - arXiv, <https://arxiv.org/html/2509.06580v1>
28. Research commissioner appears to cite discredited study in AI speech - Science|Business, <https://sciencebusiness.net/news/ai/research-commissioner-appears-cite-discredited-study-ai-speech>
29. AI is leading to thousands of job losses, report finds - CBS News, <https://www.cbsnews.com/news/ai-jobs-layoffs-us-2025/>

30. Disrupting malicious uses of AI: October 2025 | OpenAI,
<https://openai.com/global-affairs/disrupting-malicious-uses-of-ai-october-2025/>
31. AI Law, Ethics, Safety & Compliance - Cooley LLP,
<https://www.cooley.com/events/2025/2025-10-08-ai-law-ethics-safety-and-compliance>