powered by Gemini Deep Research, organized with Notebook LM, developed with chatGPT

produced by Nicole Dickens, Fractional AI Consultant

When AI Becomes Economic Infrastructure, Not Just Innovation

December 20 - December 26, 2025

Executive Summary

The final week of December 2025 will likely be recorded by future economic historians not merely as a continuation of the generative AI boom, but as the precise moment of its structural metamorphosis. We have transitioned from the era of **Generative AI**—characterized by the stochastic generation of text and images for human review—to the era of **Agentic AI**, defined by autonomous systems capable of executing end-to-end economic labor. The defining theme of the week is the "Digital Employee," a concept that has moved from theoretical whitepapers to shipping products with the release of MiniMax M2.1 and Fujitsu's Kozuchi Physical AI platform.

For the past three years, the dominant narrative has been one of "copilots"—software designed to assist a human driver. This week, the industry signaled its intent to remove the driver entirely from specific loop segments. The technological advancements released between December 20 and December 26 demonstrate a concerted effort to bridge the gap between digital reasoning and physical or systemic execution. MiniMax's new architecture allows for the direct manipulation of enterprise software interfaces, effectively treating a computer operating system not as a tool for a human, but as an environment for an agent. Simultaneously, Fujitsu's collaboration with NVIDIA has produced a "secure inter-agent gateway," a cryptographic breakthrough that allows autonomous agents from different organizations to negotiate and execute workflows without human arbitrage, promising to reduce administrative friction by nearly half.

However, this technological acceleration is colliding violently with the rigid structures of the legacy legal and political order. The week's geopolitical and regulatory news is dominated by a dual fracture. In the United States, a constitutional crisis is brewing as the executive branch moves to preempt state-level AI safety regulations via a sweeping Executive Order, pitting federal desires for "global dominance" against local mandates for safety and bias mitigation. Concurrently, the intellectual property foundation of the entire AI industry is facing its most dangerous challenge yet: a coordinated, individual lawsuit by high-profile authors, including Pulitzer Prize winner John Carreyrou, against the entire pantheon of AI giants. Unlike the class actions of 2024, this suit seeks individual statutory damages, a legal strategy that could theoretically uncap liability and force a complete "unlearning" of the datasets that power modern intelligence.

Capital markets, sensing both the opportunity of the "Agentic Turn" and the risks of the "General Model" commodity trap, have decisively shifted their attention. The "spray and pray" venture capital strategy of funding foundational LLMs is dead. In its place is a ruthless focus on **Vertical Artificial Intelligence**—specialized agents for patent law, hospital administration, and scientific discovery—and massive, sovereign-backed infrastructure plays. The smart money is no longer betting on who can build the smartest chatbot, but on who can build the most efficient "Digital Employee" for the most expensive problem.

As we close 2025, the signal to noise ratio has clarified. The novelty phase is over. We are entering a phase of industrialization, litigation, and aggressive automation. This report analyzes these critical vectors to provide a strategic roadmap for navigating the volatile transition into 2026.

Key Takeaways for SMBs: The Economics of the "Digital Employee"

The strategic landscape for Small and Medium-sized Businesses (SMBs) has fundamentally altered this week. For the last 24 months, our guidance has been consistent: adopt "Copilots" to make your existing workforce faster. That advice is now insufficient. The release of **MiniMax M2.1** and the accompanying wave of "Agentic" technologies compels us to shift our framework from "augmentation" to "autonomous execution." We are no longer discussing tools that help your accountant; we are discussing software that *is* your accounts payable department.

The Macro Shift: From Task Assistance to Workflow Ownership

The headline event of the week is the release of MiniMax M2.1, a model explicitly marketed with a "Digital Employee" mode.[1] This is not marketing hyperbole; it is a functional description of the software's capability to control mouse clicks, keyboard inputs, and navigate complex enterprise

interfaces to complete multi-step workflows without human intervention. This development matters to the SMB owner because it attacks the primary destroyer of margin in any service business: **SG&A (Selling, General, and Administrative) friction.**

In a traditional SMB, administrative workflows are fragmented. A human employee must read an invoice (step 1), verify the data against a contract (step 2), log into an ERP system (step 3), and schedule a payment (step 4). "Copilot" AI improved step 1 and perhaps step 4. The new "Digital Employee" agents handle steps 1 through 4 as a single, atomic transaction. Fujitsu's internal deployment of similar agentic technology in their purchasing department resulted in a **50% reduction in order confirmation workloads**.[2]

This effectively introduces a deflationary pressure on the cost of white-collar execution. If your competitor can execute back-office operations at half the human-capital cost by deploying autonomous agents, your existing cost structure becomes a liability. The strategic imperative is to identify "high-friction, low-judgment" workflows—such as invoice processing, compliance checking, and initial customer triage—and pilot agentic solutions immediately. The goal is not necessarily to reduce headcount, but to freeze SG&A growth while scaling revenue.

Strategic Implementation: The "Trust but Verify" Protocol

While the potential for efficiency is staggering, the operational risk is equally high. An autonomous agent that makes a mistake does not just make a typo; it can systematically corrupt a database or authorize incorrect payments at machine speed. Therefore, the implementation strategy for SMBs must be "Trust but Verify."

1. **Workflow Audit:** Map your business processes. Identify workflows that are repetitive, rules-based, and currently consume disproportionate human hours. These are your agent targets.
2. **The "Sandbox" Pilot:** Do not deploy an agentic tool on your primary ledger immediately. Run it in parallel with human execution for a quarter (a "shadow" period) to verify accuracy.
3. **Vendor Diversification:** As discussed in the *Global Policy* section below, the legal standing of major AI providers is currently under heavy litigation. Do not build your entire business logic on a single model provider (e.g., just OpenAI). Use middleware or "agent scaffolding" that allows you to swap underlying models if legal or pricing shocks occur.
4. **Service-Level Agreements (SLAs):** If you are a service provider, your clients will soon expect "agentic speed." If you take three days to turn around a legal draft or a marketing report that an agent can do in three minutes, your value proposition evaporates. You must pivot your billing model from "time and materials" to "outcomes and value."

There is a danger in over-rotating toward automation. The "Digital Employee" lacks nuance and relationship capital. In client-facing roles, the "slop" crisis (the flooding of low-quality AI content)

is creating a premium on genuine human interaction. An SMB that automates everything risks becoming indistinguishable from a spam bot. The winning strategy is to automate the back office ruthlessly while reinvesting those savings into high-touch, human-led front office service.

The Micro Signal:

This week's release of MiniMax M2.1 features a specific "Digital Employee" mode capable of operating web content via text commands, controlling mouse/keyboard inputs, and executing end-to-end tasks in finance and HR.1 This capability to interface with legacy GUI (Graphical User Interface) software bridges the gap between modern AI and older SMB software stacks.

Global AI Policy & Governance

The week of December 20, 2025, has introduced profound instability into the regulatory frameworks governing artificial intelligence. We are witnessing a collision between the accelerating capabilities of the technology and the slow-moving, yet powerful, machinery of state sovereignty and property law.

The Federal-State Constitutional Crisis in the US

A significant geopolitical fracture has opened within the United States following the mobilization of the Executive Order (EO) titled "Ensuring a National Policy Framework for Artificial Intelligence," signed on December 11, 2025.[3] This week, the operational realities of this order have begun to manifest, creating a chaotic environment for businesses operating across state lines.

The Executive Order explicitly adopts a doctrine of **Federal Preemption**, aimed at neutralizing the growing patchwork of state-level AI safety laws. States like California (with its SB 1047) and Colorado have been aggressive in legislating against AI bias and requiring safety testing. The new federal policy views these regulations not as protections, but as "onerous" obstacles to American "global AI dominance".[3]

Key Policy Mechanisms:

- **The AI Litigation Task Force:** The Department of Justice has been directed to establish a task force dedicated to challenging state AI laws in court.[3] The legal argument is likely to rest on the Commerce Clause, asserting that inconsistent state regulations burden interstate commerce in digital goods.
- **Fiscal Coercion:** The administration is leveraging the "power of the purse," threatening to withhold federal broadband funding from states that enact regulations conflicting with the "minimally burdensome" federal standard.[4]

Implications for Business:

This creates a "Schrödinger's Regulation" scenario for US-based entities. A company may find itself technically compliant with California law but arguably complicit in obstructing federal policy, or vice versa. The "safe harbor" of federal preemption is attractive for large tech firms (who lobbied for it 5), as it simplifies the compliance map to a single (lower) standard. However, the ensuing legal battles between State Attorneys General and the DOJ will likely result in years of uncertainty. SMBs and enterprises alike must prepare for a volatile compliance environment where the rules of the road are being litigated in real-time.

The "Shadow Library" Lawsuit: A Copyright Earthquake

Perhaps the most consequential governance event of the week is the copyright infringement lawsuit filed by Pulitzer Prize-winning journalist John Carreyrou (*Bad Blood*) and five other authors against the titans of the industry: **OpenAI, Google, Meta, Anthropic, xAI, and Perplexity**.[6]

Why This Lawsuit Changes Everything:

Previous legal challenges to AI training data have largely been class-action lawsuits. Class actions are often settled for "pennies on the dollar," treating the infringement as a cost of doing business (e.g., the criticized $1.5 billion Anthropic settlement 6). Carreyrou's suit is an individual action.

- **Statutory Damages:** The plaintiffs are seeking full statutory damages—up to **$150,000 per infringed work** for willful infringement.[7]
- **The "Theft" Allegation:** The complaint goes beyond simple infringement, alleging a "deliberate act of theft" through the use of "shadow libraries" (piracy hubs like LibGen and Z-Library).[7] This is a critical legal distinction. The "Fair Use" defense, which AI companies rely upon, is significantly weaker when the underlying source material was obtained illegally. Courts generally look unfavorably on "fruit of the poisonous tree."
- **Broadening the Target:** This is the first major copyright suit to name **xAI** and the search engine **Perplexity**.[7] This signals that the legal containment strategy has failed; the conflict is spreading to every player in the ecosystem, regardless of their specific architecture (e.g., Perplexity's RAG vs. OpenAI's LLM).

The "Nuclear Option" of Unlearning:

If the plaintiffs prevail, the remedy may not just be monetary. Courts could issue injunctions requiring models to be "retrained" from scratch without the infringing data—a technically arduous and exorbitantly expensive process known as "unlearning".10 This would effectively

reset the competitive landscape, erasing the advantage of companies that built their intelligence on pirated datasets.

International Divergence: EU vs. China

While the US fights internal legal battles, the international landscape is diverging.

- **European Union:** The EU is attempting to balance its heavy-handed **AI Act** with industrial reality. The European Commission's proposal to loosen environmental assessments for "AI gigafactories" [11] indicates a panic regarding competitiveness. The EU is realizing that regulation without infrastructure results in irrelevance, and is now willing to sacrifice some "Green Deal" principles to secure compute sovereignty.
- **China:** China continues to pursue a state-directed, industrial-first strategy. The funding of **DP Technology** ($114M) for "AI for Science" [12] and the release of office-automation tools like MiniMax M2.1 [1] align with Beijing's goal of using AI to upgrade *production* forces (manufacturing, science) rather than just consumer services. The Chinese ecosystem, insulated by state-backed capital and internal hedge funds, is less susceptible to the VC-driven boom-bust cycles seen in the West.[13]

While the US Executive Order aims to deregulate and accelerate AI adoption, it may paradoxically slow it down. By stripping states of their ability to regulate safety, the federal government shifts the burden of trust entirely to the tech companies. If a major AI safety incident occurs (e.g., a critical infrastructure failure caused by an autonomous agent), the lack of regulatory guardrails could lead to a massive public backlash and a subsequent regulatory overcorrection that is far more draconian than the state laws being preempted today.

AI Industry Investment

The investment activity for the week of December 20 - December 26, 2025, confirms a decisive shift in capital allocation strategies. The market has bifurcated into two distinct streams: **Strategic Infrastructure Consolidation** (Big Tech locking down the physical supply chain) and **Vertical SaaS** (VCs betting on specific, high-value workflows).

1. The Infrastructure Supercycle: $24.75 Billion in Consolidation

While venture capital chased software, the industry titans executed nearly **$25 billion** in deals this week to secure the physical inputs of AI: chips and energy. This is not investment; it is industrial capture.

- **Nvidia Acquires Groq ($20B):** In the single largest move of the week, Nvidia agreed to acquire assets from AI chip startup Groq in an all-cash deal valued at **$20 billion**.

    - **The Strategy:** Though technically structured as a non-exclusive licensing agreement and talent acquisition ("acqui-hire"), the deal effectively neutralizes a potent competitor in the *inference* market. Groq's "Language Processing Unit" (LPU) technology was the fastest on the market for running models. By absorbing Groq's founder Jonathan Ross and his engineering team, Nvidia is moving to monopolize not just the training of AI (H100/Blackwell chips), but the *running* of AI (Inference).
    - **Implication:** This signals that the "inference cost race to the zero" may slow down. If Nvidia controls the best inference tech, they control the margins of every AI application startup.
- **Alphabet (Google) Acquires Intersect Power ($4.75B):** Google parent Alphabet signed a definitive agreement to acquire energy infrastructure firm Intersect Power for **$4.75 billion**.
    - **The Strategy:** This is a vertical integration of the power grid. As AI data centers consume gigawatts of power, electricity availability has become the primary bottleneck for growth. Google is no longer waiting for utility companies to build capacity; they are buying the power developer directly to secure "multiple gigawatts" of renewable energy for their data centers.
    - **Implication:** We are entering an era where Big Tech companies are effectively becoming energy utilities. Access to clean, massive-scale power is now a proprietary competitive advantage, not a public commodity.

2. The Rise of Vertical Agentic AI

Below the mega-cap infrastructure layer, smart venture capital is flowing into startups that own specific vertical workflows.

Table 1: Key Venture Capital Rounds (Dec 20 - Dec 26, 2025)

Company	Sector	Amount Raised	Lead Investor	Strategic Focus
DP Technology	Scientific AI	$114M (800M RMB)	Fortune VC, Beijing Jingguorui	"AI for Science" platform (drug discovery, molecular simulation).[12]
Ankar	Legal Tech	$20M (Series A)	Atomico	Automating patent drafting and prosecution using AI agents.[14]
Mindoo	Health Tech	€5M (Seed)	Pitchdrive, QBIC	AI agents for reducing hospital

				administrative workload.[14]
Tin Can	Consumer HW	$12M (Seed)	Greylock Partners	"Landline-style" phone for kids (tech-minimalism trend).[16]
Nawah Scientific	Deep Tech	$23M (Series A)	Life Ventures	Scientific research enablement (Egypt-based).[16]

Analysis of Verticalization:

- **Ankar's Patent Agents:** The $20M raise for Ankar is emblematic. Patent law is a text-heavy, logic-heavy, and high-cost discipline. By building specialized agents trained on 150 million patent applications [15], Ankar is not just "helping" lawyers; it is industrializing the production of intellectual property.
- **Scientific Discovery as a Service:** DP Technology's massive $114M round highlights the growing belief that the highest ROI for AI is not in writing emails, but in simulating physics. Their "Particle Universe" simulation system [12] represents the convergence of High-Performance Computing (HPC) and AI, a trend echoed by Fujitsu's roadmap.

Consolidation and the "Service-as-Software" Pivot

We are also witnessing a wave of consolidation where traditional service providers acquire technology to transform their business models.

- **1Point1 Solutions Acquires Netcom ($33.37M):** This acquisition is not merely about geographic expansion into Latin America. The stated strategy is to deploy "GenAI-enabled platforms" across Netcom's delivery centers.[17] This is the "Service-as-Software" thesis in action: buying a traditional labor-intensive BPO (Business Process Outsourcing) firm and upgrading its margins by replacing human labor with AI agents. This validates the "Digital Employee" macro trend from an M&A perspective.
- **The "Quasi-Merger" Crackdown:** Regulatory pressure is forcing big tech to be creative with M&A. The "acqui-hire" model—seen vividly in the **Nvidia/Groq** deal where the team moves but the legal entity technically remains independent—is a direct response to antitrust scrutiny. This allows giants to acquire capability without triggering traditional merger reviews.[19]

There is a "Capital Indigestion" risk. The massive inflows into infrastructure (Nvidia/Google deals totaling ~$25B) are predicated on the assumption of massive future revenue. However, reports suggest that many enterprises are still in "pilot purgatory," struggling to prove ROI for

generative AI.20 If the "application layer" (like Ankar or Mindoo) does not scale fast enough to consume the massive compute capacity being built, we could face a "fiber optic bubble" scenario similar to 2000—a crash in infrastructure prices due to oversupply.

Breakthroughs in AI Technology

The technological advancements of late December 2025 are less about raw model size and more about **architecture** and **interoperability**. The industry is solving the engineering challenges required to make "Agentic AI" reliable, affordable, and capable of interacting with the physical world.

1. MiniMax M2.1: The Economics of MoE Agents

The release of **MiniMax M2.1** [1] represents a maturation in model architecture designed for economic viability.

- **Mixture-of-Experts (MoE):** The model utilizes a massive 230B parameter count but only activates **10B parameters per token** generation.[22] This "sparse" activation drastically reduces the computational cost of running the model. For an SMB deploying a "Digital Employee" that runs 24/7, this inference cost efficiency is the difference between a profitable tool and a money pit.
- **Agent Scaffolding:** The innovation is not just the weights, but the "scaffolding"—the software wrapper that allows the model to manage long-context memories, execute multi-step reasoning chains, and interface with external tools.[1]
- **VIBE Benchmark:** MiniMax introduced the "VIBE" benchmark to test "interactive logic and visual aesthetics," scoring highly on generating functional apps.[1] This signals a shift in evaluation metrics from "academic benchmarks" (like MMLU) to "production benchmarks" (can it build a working app?).

2. Fujitsu Kozuchi Physical AI 1.0: The Secure Bridge

Fujitsu's collaboration with NVIDIA to release **Kozuchi Physical AI 1.0** [2] addresses two critical bottlenecks: **Security** and **Physicality**.

- **Secure Inter-Agent Gateway:** As agents proliferate, they need to talk to each other (e.g., a supplier's agent negotiating with a buyer's agent). The "Secure Inter-Agent Gateway" uses cryptographic protocols to allow these agents to collaborate without exposing the underlying confidential data of either party.[2] This is a prerequisite for B2B agentic commerce.
- **Physical AI:** The platform extends agentic control to robotics. In Fujitsu's vision, an agent doesn't just process a procurement order; it directs a physical robot in a warehouse to

execute the logistics. This moves AI from the screen to the factory floor.

3. NVIDIA NIM Microservices: The Deployment Standard

Underpinning these advances is the standardization of deployment via **NVIDIA NIM (Inference Microservices)**.[24]

- **Containerized Intelligence:** NIMs provide a standardized way to package and deploy models like Mistral or Llama on any NVIDIA-accelerated infrastructure. This week's updates optimized these containers for agentic workflows, effectively creating a "Docker for AI Agents." This lowers the technical barrier for enterprises to host their own sovereign agents, reducing dependency on public APIs.

Table 2: Technological Comparison of Week's Releases

Feature	MiniMax M2.1	Fujitsu Kozuchi Physical AI
Primary Focus	Digital Office Automation ("Digital Employee")	Physical/Robotic Integration & Security
Architecture	Mixture-of-Experts (MoE) 230B/10B Active	Agentic Framework + NVIDIA Stack
Key Capability	End-to-end GUI control (Mouse/Keyboard)	Secure Inter-Agent Negotiation
Target Use Case	Finance, HR, Coding, Admin	Supply Chain, Manufacturing, Procurement
Deployment	API & Open Weights	Enterprise Platform (On-prem/Cloud)

The "Digital Employee" concept relies heavily on the reliability of the underlying model. While MiniMax M2.1 performs well on benchmarks, "cringe" marketing videos and early user feedback suggest that the user experience (UX) and prompt adherence can still be finicky compared to established leaders like Gemini or GPT-4. The "last mile" of reliability—where an agent works 99.9% of the time—remains the hardest engineering challenge.

Societal and Economic Implications

The convergence of "Digital Employees," massive capital concentration, and legal warfare creates a volatile societal mixture. We are moving from a period of "AI Hype" to "AI Reality," and the friction of this transition is generating heat in the labor market and the culture.

The "White Collar" Labor Disruption

The branding of MiniMax M2.1 as a "Digital Employee" is a semantic milestone. It explicitly frames the technology as a labor substitute.

- **Deflationary Shock:** If Fujitsu's internal data holds true—that agentic AI can reduce administrative workload by **50%** [2]—we are facing a significant contraction in demand for entry-level white-collar labor. The roles that typically serve as "apprenticeships" for the middle class (data entry, junior compliance, basic coding) are the exact targets of these new agents.
- **The "Hollow Middle":** The danger is a "hollow middle" corporate structure, where you have senior decision-makers (who direct the agents) and the agents themselves, with very few rungs on the ladder for junior employees to climb. This creates a long-term crisis in human capital development: how do you train a Senior Strategist if no one is hiring Junior Analysts?

The "Slop" Crisis and the Epistemology of the Web

The proliferation of agentic content generation is degrading the quality of the open web, a phenomenon increasingly referred to as "Slop."

- **Automated Propaganda:** The controversies surrounding Elon Musk's **Grok** chatbot—generating far-right misinformation and "hallucinated" entries in its "Grokipedia" feature [5]—demonstrate the risks of uncurated agentic output. When agents can generate content at infinite scale, the "public square" risks being flooded with synthetic noise.
- **The Trust Premium:** As "slop" proliferates, the value of *verified human authentication* increases. We may see a bifurcated internet: a "High Trust" web of verified human users (paid, gated) and a "Low Trust" open web dominated by agents screaming at agents.

The "Bubble" Narrative and the Human Correction

There is a growing intellectual counter-narrative, articulated by Rafael Behr, that the AI industry is an "Icarus Economy" or a financial bubble.[5]

- **The Correction as Opportunity:** Behr argues that the current valuation bubble is "puffed up" by Silicon Valley ego and unsustainable CapEx. He suggests that the inevitable "bursting" of this bubble is not a disaster, but a necessary correction. It is only when the "money hose" turns off that society will have the leverage to impose meaningful democratic control and regulation over these technologies. The Carreyrou lawsuit may be the pin that pricks this bubble, forcing the industry to reckon with the true cost of its raw materials.

While the "Labor Displacement" narrative is strong, history suggests that technology often creates new, unforeseen categories of work. The "Prompt Engineer" role of 2023 is already

evolving into the "Agent Orchestrator" role of 2026. The demand for human judgment, empathy, and creative strategy may actually increase as the cost of execution falls to zero. The "Digital Employee" may not fire the human; it may simply force the human to be more human.

Sources

- [11] Jisc, "December 2025 round-up of interesting AI news"
- [23] Fujitsu, "Fujitsu develops Fujitsu Kozuchi Physical AI 1.0"
- [24] NVIDIA, "Nemotron 3 Nano and Mistral 3 Updates"
- [4] Sidley Austin, "Unpacking the December 11, 2025 Executive Order"
- [27] Mayer Brown, "President Trump Issues Executive Order on AI Policy"
- [3] White House, "Executive Order: Eliminating State Law Obstruction"
- [5] The Guardian, "When the AI bubble bursts..." (Rafael Behr)
- [6] Irish Sun, "New York Times reporter sues Google, xAI, OpenAI"
- [14] Crescendo.ai, "Ankar Raises $20M Series A"
- [28] American Bazaar, "10 Biggest AI Deals of 2025"
- [17] TechCircle, "1Point1 Solutions to acquire Netcom"
- [7] The Decoder, "Authors sue six AI giants for book piracy"
- [1] MiniMax, "MiniMax M2.1 Release Notes"
- [21] TechRepublic, "MiniMax Unveils M2.1"
- [22] Microsoft Azure Blog, "MiniMax M2 in Azure"
- [2] TechLoy, "Fujitsu rolls out Kozuchi Physical AI 1.0"
- [29] Hyperight, "Fujitsu launches Kozuchi Physical AI 1.0"
- [30] Crunchbase, "Big Funding Trends Charts EOY 2025"
- [14] Crescendo.ai, "Latest VC Investment Deals"
- [15] Index Ventures, "Ankar Raises $20M"
- [18] CXOToday, "1Point1 Solutions Acquisition"
- [13] American Bazaar, "China unveils national venture capital fund"
- [16] PitchBook, "Global VC Deals Dec 23, 2025"
- [31] SWF Institute, "Sovereigns of Silicon Valley"
- [12] SiliconANGLE, "DP Technology raises $114M"
- [20] IndexBox, "VCs Raise the Bar for Startup Funding"
- [19] MarketMinute, "The Quasi-Merger Crackdown"
- [26] Hacker News, "Reaction to MiniMax M2.1"
- [32] eWeek, "Pulitzer Winner Sues 6 AI Giants"
- [9] Modern Diplomacy, "John Carreyrou exposes a new fraud"
- [10] i10x.ai, "AI Copyright Lawsuit: What happened"
- [8] eWeek, "Implications of Lawsuit"
- [5] The Guardian, "Arguments in 'When the AI bubble bursts'"

- MarketMinute, "Nvidia's $20 Billion Christmas Eve Coup"
- The Tech Buzz, "Nvidia Snaps Up Groq in Record $20B Deal"
- Groq Newsroom, "Groq and Nvidia Enter Licensing Agreement"
- TechInformed, "Alphabet to Buy Intersect for Nearly $5B"
- Alphabet Investor Relations, "Alphabet Announces Agreement to Acquire Intersect"
- Times of India, "Google Makes $4.75 Billion Acquisition"

Works cited

1. MiniMax M2.1: Significantly Enhanced Multi-Language Programming, Built for Real-World Complex Tasks, accessed December 26, 2025, https://www.minimax.io/news/minimax-m21
2. Fujitsu rolls out Kozuchi Physical AI 1.0 with NVIDIA to power agentic automation, accessed December 26, 2025, https://www.techloy.com/fujitsu-rolls-out-kozuchi-physical-ai-1-0-with-nvidia-to-power-agentic-automation/
3. Ensuring a National Policy Framework for Artificial Intelligence - The White House, accessed December 26, 2025, https://www.whitehouse.gov/presidential-actions/2025/12/eliminating-state-law-obstruction-of-national-artificial-intelligence-policy/
4. Unpacking the December 11, 2025 Executive Order: Ensuring a National Policy Framework for Artificial Intelligence, accessed December 26, 2025, https://www.sidley.com/en/insights/newsupdates/2025/12/unpacking-the-december-11-2025-executive-order
5. When the AI bubble bursts, humans will finally have their chance to ..., accessed December 26, 2025, https://www.theguardian.com/commentisfree/2025/dec/23/artificial-intelligence-ai-bubble-bursts-humans-take-back-control
6. Google, xAI, OpenAI sued by NYT reporter for using copyrighted book, accessed December 26, 2025, https://www.irishsun.com/news/278771765/google-xai-openai-sued-by-nyt-reporter-for-using-copyrighted-book
7. Authors sue six AI giants for book piracy - The Decoder, accessed December 26, 2025, https://the-decoder.com/authors-sue-six-ai-giants-for-book-piracy/
8. Pulitzer Winner Sues 6 AI Giants for Billions Over Copyright - eWeek, accessed December 26, 2025, https://www.eweek.com/news/pulitzer-winner-sues-ai-giants/
9. John Carreyrou Exposes a New Fraud: The Theft Behind Your Chatbot - Modern Diplomacy, accessed December 26, 2025, https://moderndiplomacy.eu/2025/12/23/john-carreyrou-exposes-a-new-fraud-the-theft-behind-your-chatbot/
10. AI Copyright Lawsuit Hits OpenAI, Google, Meta & More - i10X, accessed December 26, 2025, https://i10x.ai/ru/news/ai-copyright-lawsuit-authors-sue-openai-google
11. December 2025 round-up of interesting AI news and announcements - Artificial intelligence, accessed December 26, 2025,

https://nationalcentreforai.jiscinvolve.org/wp/2025/12/18/december-2025-round-up-of-interesting-ai-news-and-announcements/
12. DP Technology raises $114M to accelerate China's AI for science industry, accessed December 26, 2025, https://siliconangle.com/2025/12/24/dp-technology-raises-114m-accelerate-chinas-ai-science-industry/
13. China unveils national venture capital fund to boost innovation, accessed December 26, 2025, https://americanbazaaronline.com/2025/12/26/china-unveils-national-venture-capital-fund-to-boost-innovation-472222/
14. Latest AI Startup Funding News and VC Investment Deals - 2025 - Crescendo.ai, accessed December 26, 2025, https://www.crescendo.ai/news/latest-vc-investment-deals-in-ai-startups
15. Ankar Raises $20M Series A to Accelerate and Protect Innovation - Index Ventures, accessed December 26, 2025, https://www.indexventures.com/perspectives/ankar-raises-20m-series-a-to-accelerate-and-protect-innovation/
16. Global VC deals for December 23, 2025 - PitchBook, accessed December 26, 2025, https://pitchbook.com/newsletter/global-vc-deals-for-december-23-2025
17. It's a wrap: News this week (December 20-26) - Techcircle, accessed December 26, 2025, https://www.techcircle.in/2025/12/26/it-s-a-wrap-news-this-week-december-20-26
18. 1Point1 Solutions' USD 33.37 Million Netcom Acquisition Marks Entry into Latin America; Set to Nearly Double FY27 Revenue - CXOToday.com, accessed December 26, 2025, https://cxotoday.com/press-release/1point1-solutions-acquires-netcom-expands-global-footprint-to-latin-america/
19. The "Quasi-Merger" Crackdown: How US Regulators are Redrawing the Map for Big Tech's AI Ambitions, accessed December 26, 2025, https://markets.financialcontent.com/stocks/article/marketminute-2025-12-26-the-quasi-merger-crackdown-how-us-regulators-are-redrawing-the-map-for-big-techs-ai-ambitions
20. Venture Capital 2026: Higher Funding Bar, Focus on Proof & Momentum - News and Statistics, accessed December 26, 2025, https://www.indexbox.io/blog/vcs-raise-the-bar-for-startup-funding-in-2026-shift-from-vision-to-proof/
21. MiniMax Unveils M2.1 to Bring Multilingual Programming Gains to Open AI Models, accessed December 26, 2025, https://www.techrepublic.com/article/news-minimax-m2-release/
22. MiniMax-M2: The Open-Source Innovator in Coding and Agentic Workflows Now in Azure AI Foundry | Microsoft Community Hub, accessed December 26, 2025, https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/minimax-m2-the-open-source-innovator-in-coding-and-agentic-workflows-now-in-azur/4466045
23. Fujitsu develops Fujitsu Kozuchi Physical AI 1.0 for seamless ..., accessed

December 26, 2025, https://global.fujitsu/en-global/pr/news/2025/12/24-02
24. ICYMI: Top AI Model Releases - December 2025 - YouTube, accessed December 26, 2025, https://www.youtube.com/shorts/P9KIN1iQAC8
25. Artificial Intelligence 2025 Legislation - National Conference of State Legislatures, accessed December 26, 2025, https://www.ncsl.org/technology-and-communication/artificial-intelligence-2025-legislation
26. MiniMax M2.1: Built for Real-World Complex Tasks, Multi-Language Programming, accessed December 26, 2025, https://news.ycombinator.com/item?id=46388213
27. President Trump Issues Executive Order on "Ensuring a National Policy Framework for Artificial Intelligence" | Insights | Mayer Brown, accessed December 26, 2025, https://www.mayerbrown.com/en/insights/publications/2025/12/president-trump-issues-executive-order-on-ensuring-a-national-policy-framework-for-artificial-intelligence
28. 10 biggest AI deals of 2025: The year AI money went nuclear, accessed December 26, 2025, https://americanbazaaronline.com/2025/12/24/10-biggest-ai-deals-of-2025-the-year-ai-money-went-nuclear-472113/
29. Fujitsu Launches Kozuchi Physical AI 1.0 to Unify Robotics and Agents - hyperight.com, accessed December 26, 2025, https://hyperight.com/fujitsu-launches-kozuchi-physical-ai-1-0-to-unify-robotics-and-agents/
30. 6 Charts That Show The Big AI Funding Trends Of 2025 - Crunchbase News, accessed December 26, 2025, https://news.crunchbase.com/ai/big-funding-trends-charts-eoy-2025/
31. Meet the Sovereigns of Silicon Valley: How Nation-States Are Bankrolling the AI Revolution, accessed December 26, 2025, https://www.swfinstitute.org/news/108142/meet-the-sovereigns-of-silicon-valley-how-nation-states-are-bankrolling-the-ai-revolution
32. Pulitzer Winner Sues 6 AI Giants for Billions Over Copyright - eWeek, accessed December 26, 2025, http://www.eweek.com/news/pulitzer-winner-sues-ai-giants/